

• OMNIS

the partly omniscient auditor

Automated Compliance Evidence Collection & Audit
Société Générale · a compliance evidence engine

OMNISCIENCE INDEX

92.4 /100

synthetic bench · runs offline

You can be compliant and still fail the audit.

The control works. The proof is scattered — one config in AWS, one screenshot saved six months ago, one that technically exists but nobody can find. A control fails the audit because no one located the receipt.

72_h


HUNTING EVIDENCE, EVERY CYCLE

OMNIS reports one honest number — and names what it can't prove.

$$\begin{array}{ccccc} \text{COVERAGE} & & \text{FRESHNESS} & & \text{CONFIDENCE} \\ \text{w } 0.40 & \times & \text{w } 0.30 & \times & \text{w } 0.30 \end{array} = \begin{array}{c} \text{OMNISCIENCE INDEX} \\ 0-100 \end{array}$$

A requirement with no evidence is reported **UNKNOWN** — counted against the score, never quietly skipped. No auditor is omniscient; this one tells you exactly where the gaps are.

Files in, an auditor-ready report out — one offline pipeline.



Each box is a real module with its own test suite. Modular, offline, no API key — and an optional LLM adapter that ships off, never on the critical path.

Every link shows its method; every hard case keeps its name.

THE LAYERED LINKER

1	exact id	cheapest, surest
2	framework rule	type + standard
3	TF-IDF	offline, on text
4	UNMAPPED	honest, not forced

THE FOUR HARD CASES

missing	→	UNKNOWN
conflicting	→	PARTIAL
low-confidence	→	caught by integrity
ambiguous	→	flagged at parse

p = 0.86

permutation test · best real precision 0.333

**We audited the benchmark
before trusting it. The
labels are noise.**

*"We built a compliance auditor. The first thing it audited was
the benchmark."*

01 SHIPPED

a reproducible finding
make analyze

02 BUILT

an honest synthetic bench
labels by construction

03 WIRED

a re-grade flag
--labels, one command

Every bar the brief set, cleared with room to spare.

.957 / .953

DETECTION P / R
bar 0.70 / 0.60

78%

AUTOMATION RATE
clears the 70% target

< 1s

FULL PIPELINE
500 + 5,000 · bar 60s

54.7%

RECOVERY, ID HIDDEN
8× the 6.7% random

92.4

OMNISCIENCE INDEX
out of 100 · synthetic

102

TESTS PASSING
one suite per module

It leaves the gaps visible instead of hiding them.

Real evidence corpora are ugly. OMNIS reports the defects, with a severity, rather than crashing on them.

HIGH	500	rows share one identical description
HIGH	500	cite requirement IDs in no policy (381 unique)
MED	239	reviewed before collected – impossible
MED	455	stored freshness disagrees with dates
MED	18	Approved but confidence < 0.6

UNKNOWN = missing evidence UNMAPPED = fits nothing

One page an auditor reads; one report an auditor accepts.

OMNIS

the partly omniscient auditor

Provided sample

Synthetic

OMNISCIENCE INDEX

92.6 / 100

How much we know, and how sure we are: a weighted blend of evidence coverage, freshness, and confidence across every requirement.

64.4%

AUTOMATION RATE

500

EVIDENCE RECORDS

STATUS DISTRIBUTION

12 partial

3 gap

ASK

Try: show gaps, encryption, unknown, logging

show gaps

encryption

unknown

logging

REQUIREMENTS

REQUIREMENT	STATUS	CONFIDENCE	EVIDENCE	RATIONALE
POL-ENC-001-R1 All data at rest must be encrypted using AES-256 or stronger	PARTIAL	0.33	29	Mixed signals: 9 good, 6 failing, 14 ambiguous record(s).
POL-ENC-001-R2 Encryption keys must be rotated at least annually	PARTIAL	0.51	31	Mixed signals: 17 good, 7 failing, 7 ambiguous record(s).
POL-ENC-001-R3 Encryption keys must be stored securely	PARTIAL	0.33	10	10 record(s) affirmatively failing (mixed /

make demo · single-page dashboard, stdlib server

OMNIS Compliance Evidence Report

the partly omniscient auditor

synthetic bench · 15 requirements · generated 2026-04-15

Abstract

OMNIS parsed 15 requirements from the synthetic policy set and linked 500 evidence records against them. The Omniscience Index, a weighted composite of evidence coverage, freshness, and confidence, is 92.4 out of 100, with an automation rate of 78.0%. 29 evidence records were left unmapped. The sections below give the per-requirement compliance status with its supporting evidence and freshness, followed by the data-quality findings OMNIS raised against the corpus itself.

1 Executive summary

Metric	Value
--------	-------

make report · auditor-ready PDF, typeset

An auditor accepts it because it shows its work.

-
- | | | |
|----|--------------------|--|
| 01 | It shows its work | every link carries its method and a confidence; every verdict, its evidence. |
| 02 | It is reproducible | offline, in one command, no API key — the same result every run. |
| 03 | It is honest | it states what it cannot prove instead of guessing. |
-

• OMNIS

github.com/ataylus/OMNIS

make demo • make test • make analyze
offline • no API key

reports/report.pdf • docs/detailed
the output • the full writeup

No auditor is omniscient. This one tells you
exactly where the gaps are.